

# Linux cgroups

# Contents

---

`cgroups -$ ./tutorial`

- (1) Introduction
- (2) Using cgroups

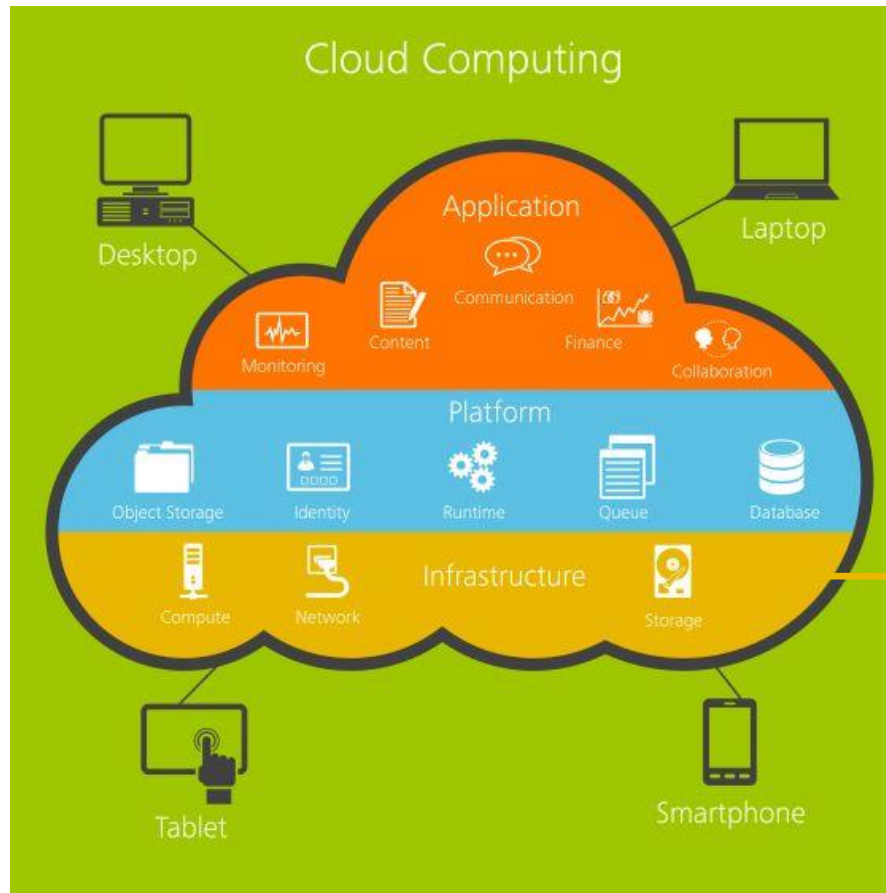
> 1

Introduction

- | - organization
- | - tasks and subsystems
- | - hierarchies and rules


> \_

# Whetting your appetite




App specs:

- CPU time
- System memory
- Disk bandwidth
- Network bandwidth
- Monitoring



Feedback



Monitor



- CPU?
- Memory?
- iops/bytes in/out?

Group subsys requirements:

1. 60% CPU
2. 14 GB memory
3. 60% r/w blk dev
4. 80 net bw

Analyze



# What are cgroups?

---

"Control Groups provide a mechanism for aggregating/partitioning sets of **tasks**, and **all their future children**, into hierarchical groups described by specialized behavior. "

- lxr

# What are control groups about?

What

Resource allocation management:

- CPU time
- System memory
- IOPS
- Network bandwidth

demanded by a group of tasks (processes) → **cgroup**

How

Operations – functionality:

- ✓ Monitoring
- ✓ Resource access denial
- ✓ Cgroup reconfiguration on-the-fly

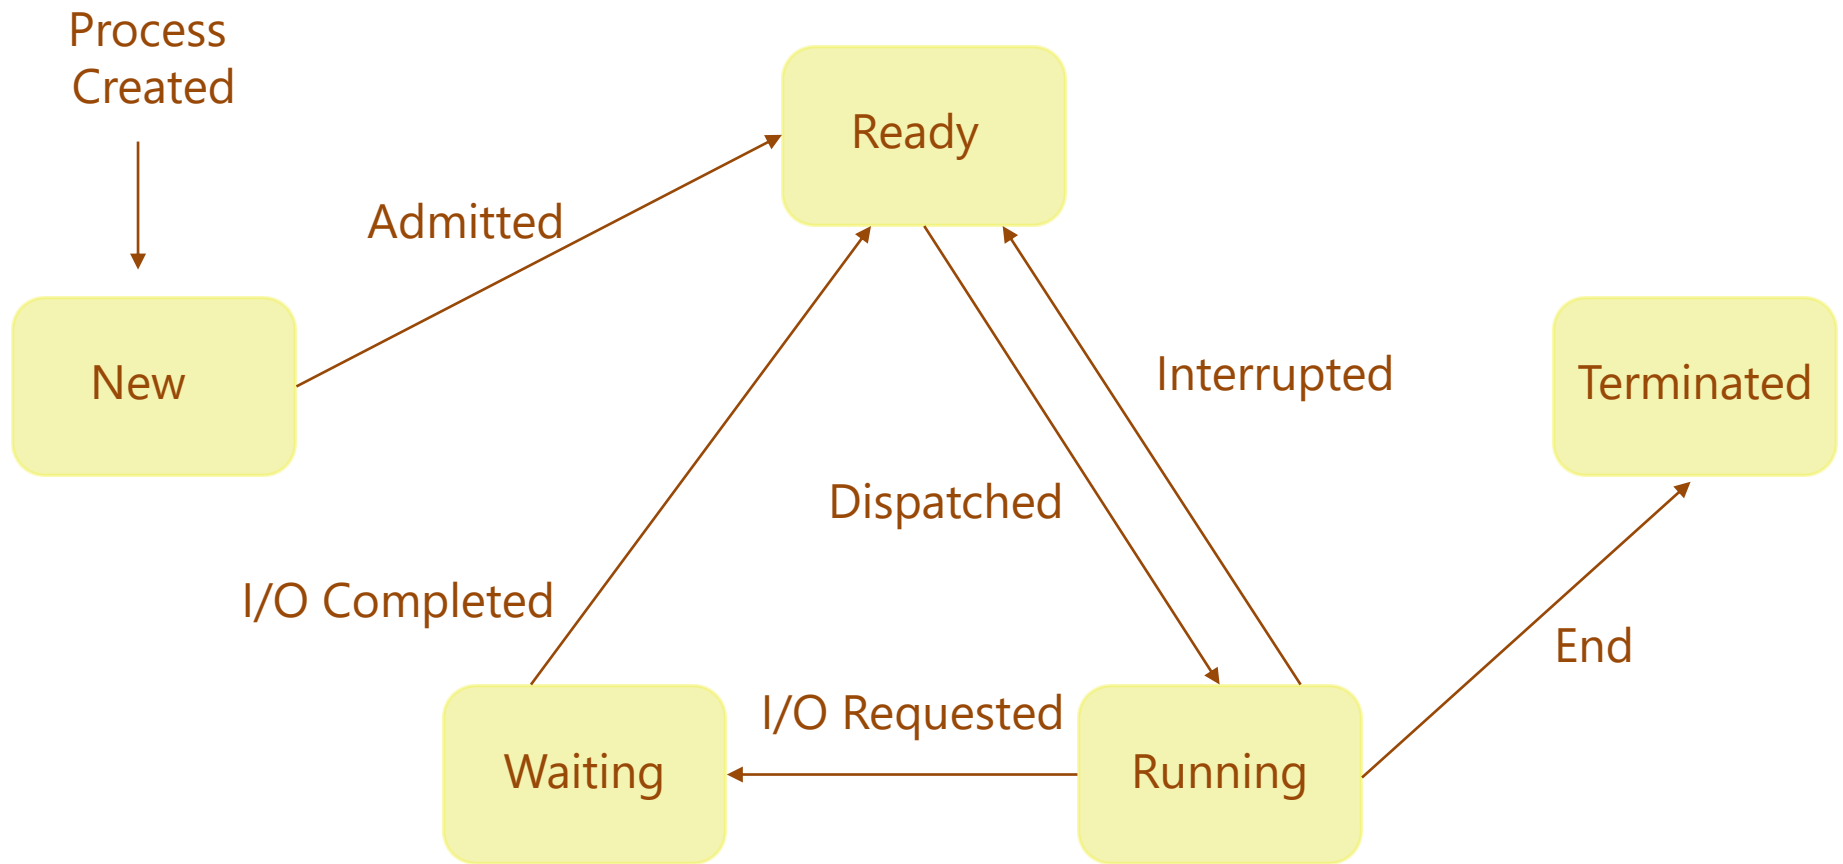
Why

Fined grained control over:

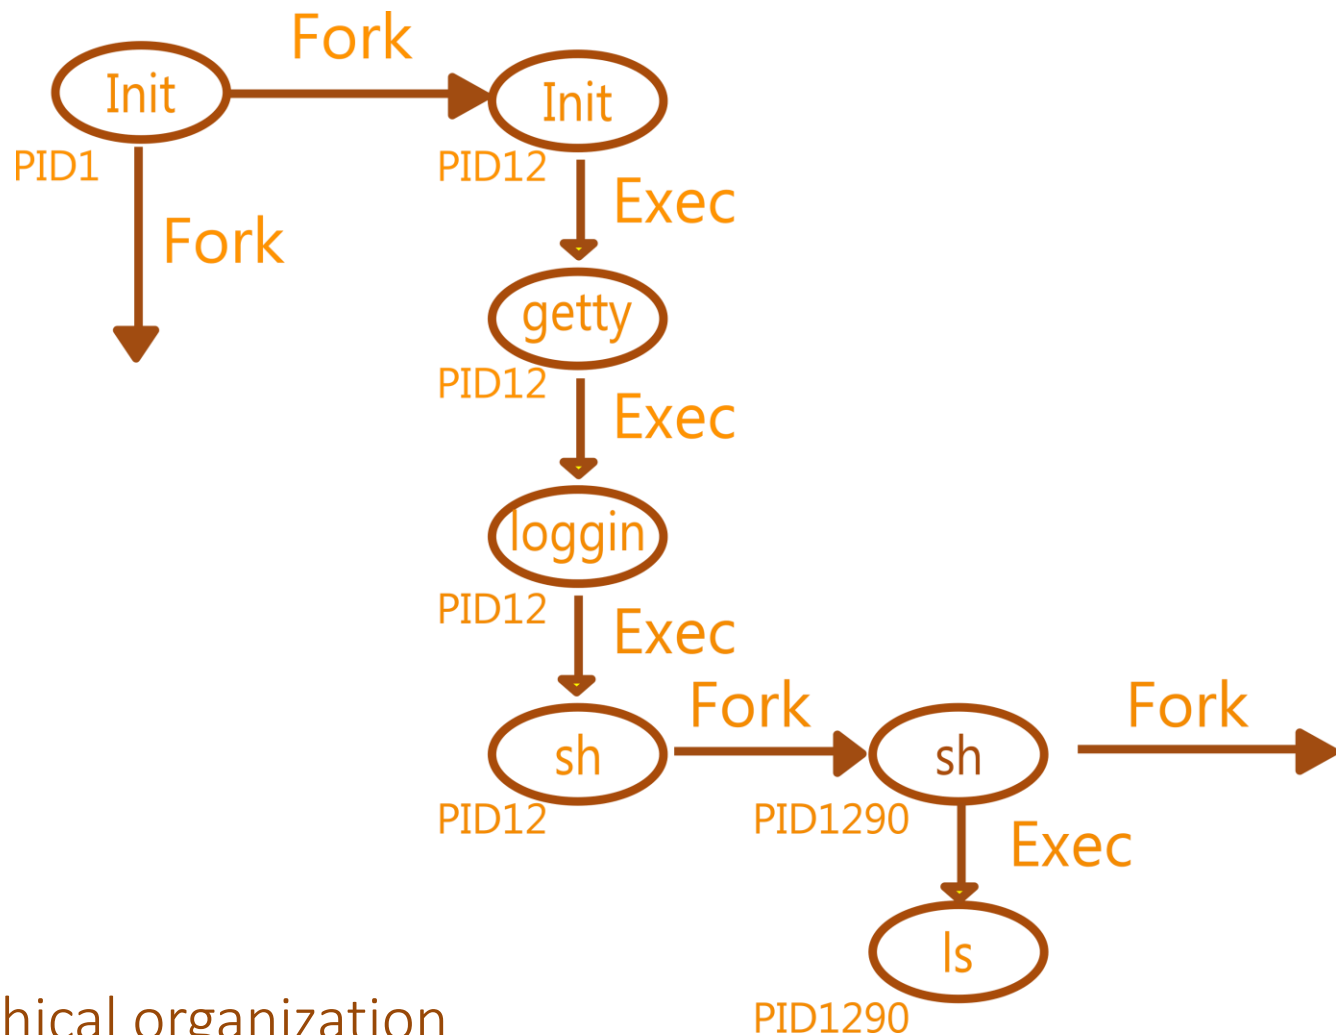
allocating, prioritizing, denying, managing and monitoring  
system resources to provide

overall system **efficiency**

# Let's recall the Linux Process Model!







- > Hierarchical organization
- > Attribute inheritance

# Control group organization

---

“Cgroup ::= user defined group of processes attached to user defined resource management plans/policies”

- Hierarchical organization
  - Consider linux processes organization → tree
  - Cgroups organization → forest (with trees sharing leaves!)

- Attribute inheritance

+ many different **hierarchies** of cgroups can exist simultaneously

+ each hierarchy is **attached** to one or more **subsystems**

# Subsystem ?

---

“ A \*subsystem\* is a module that makes use of the task grouping facilities provided by cgroups to treat groups of tasks in particular ways.

A subsystem is typically a "**resource controller**" that **schedules** a resource or applies per-cgroup **limits**, but it may be **anything that wants to act on a group of processes**, e.g. a virtualization subsystem. ”

- lxr

# Available to cgroups subsystems

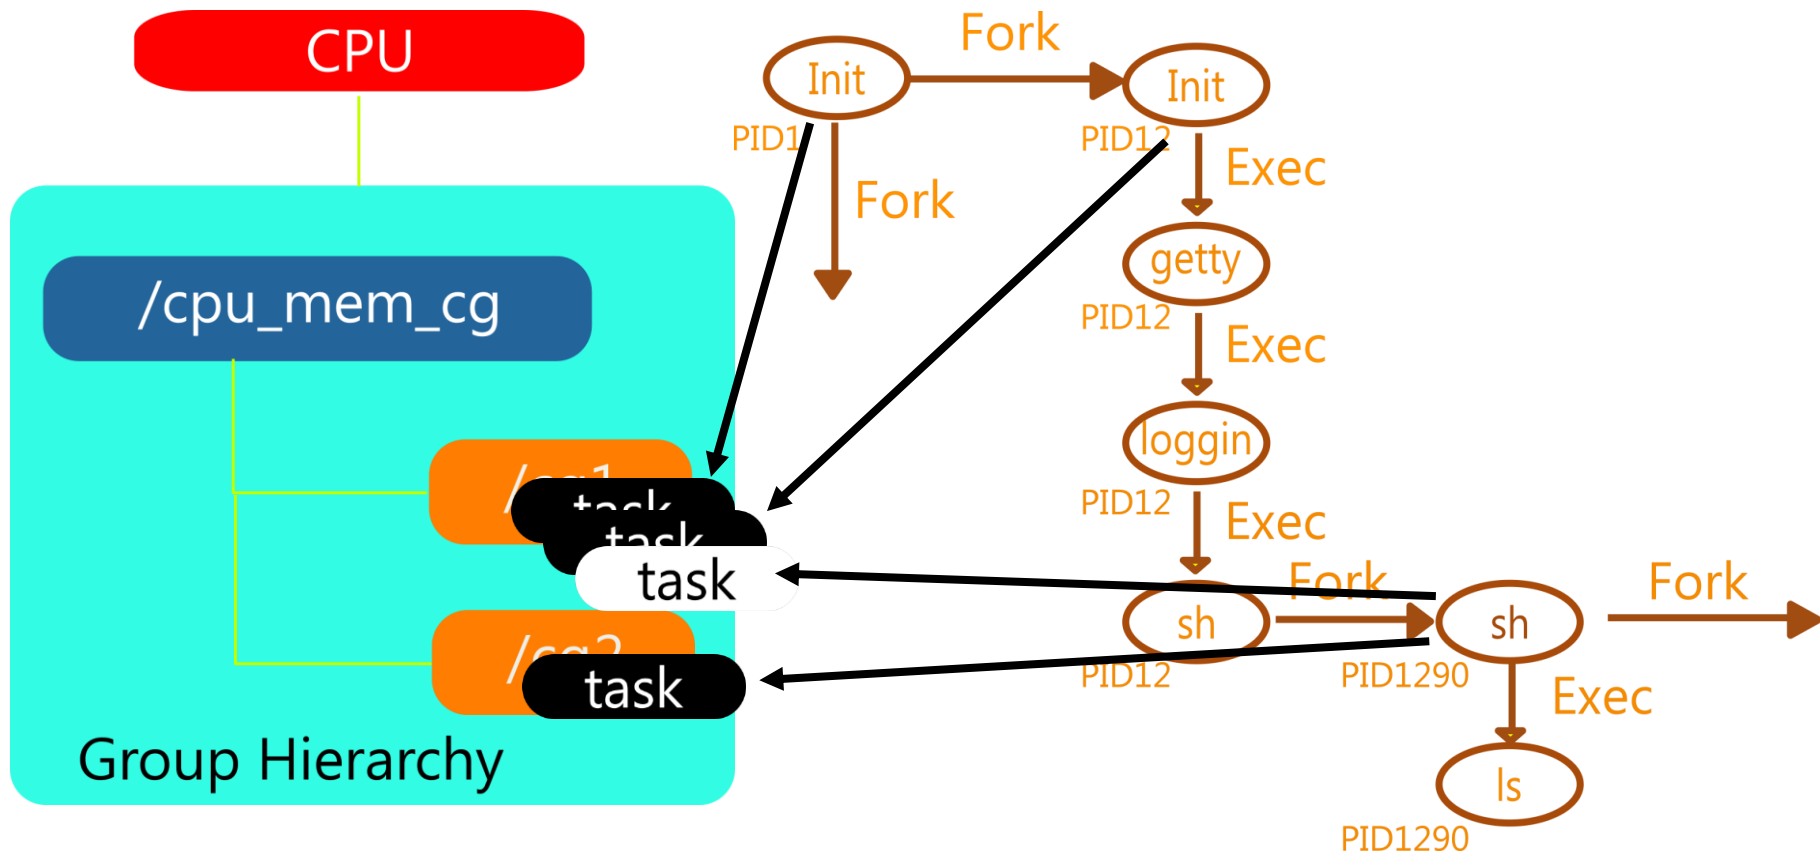
- **blkio** — this subsystem sets limits on input/output access to and from block devices such as physical drives (disk, solid state, USB, etc.)
- **cpu** — this subsystem uses the scheduler to provide cgroup tasks access to the CPU
- **cpuacct** — this subsystem generates automatic reports on CPU resources used by tasks in a cgroup
- **cpuset** — this subsystem assigns individual CPUs (on a multicore system) and memory nodes to tasks in a cgroup
- **devices** — this subsystem allows or denies access to devices by tasks in a cgroup
- **freezer** — this subsystem suspends or resumes tasks in a cgroup
- **memory** — this subsystem sets limits on memory use by tasks in a cgroup, and generates automatic reports on memory resources used by those tasks
- **net\_cls** — this subsystem tags network packets with a class identifier (classid) that allows the Linux traffic controller (tc) to identify packets originating from a particular cgroup task
- **net\_prio** — this subsystem provides a way to dynamically set the priority of network traffic per network interface
- **ns** — the namespace subsystem \*

# Hierarchy?

---

A \*hierarchy\* is a set of cgroups arranged in a tree, such that every task in the system is in exactly one of the cgroups in the hierarchy, and a set of subsystems; each subsystem has system-specific state attached to each cgroup in the hierarchy. Each hierarchy has an instance of the cgroup virtual filesystem associated with it.

- lxr

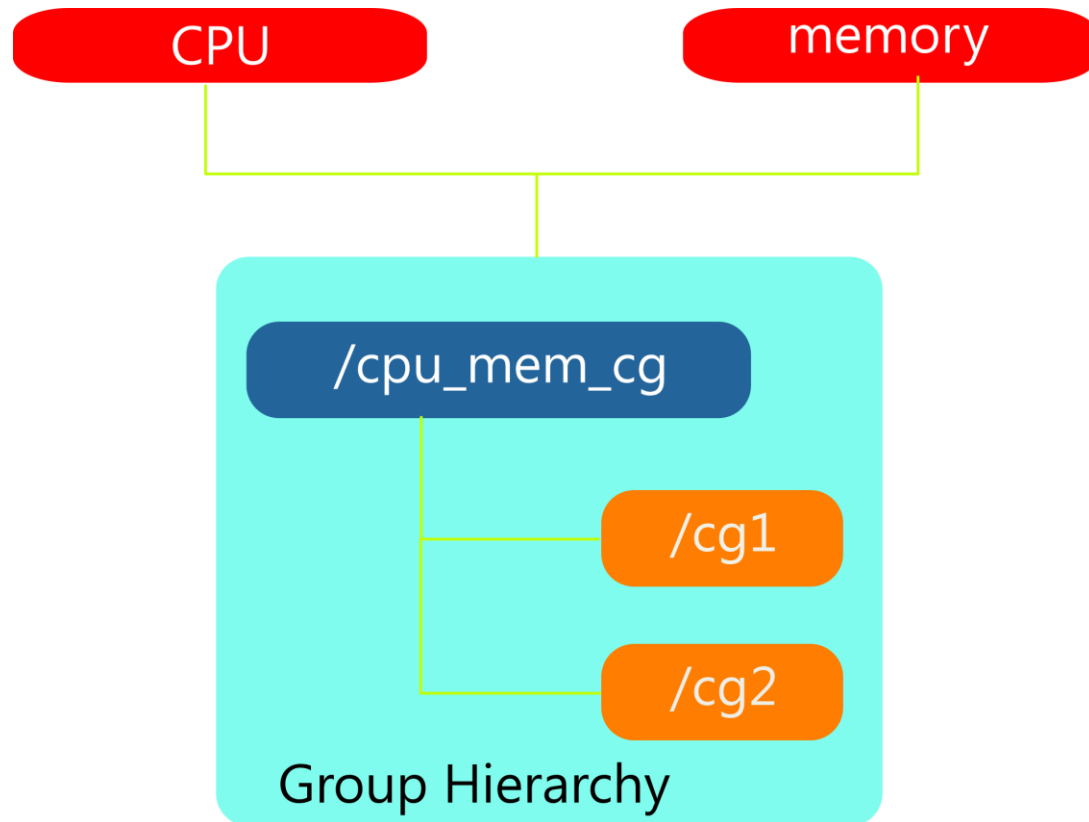


# Attaching tasks to subsystems via cgroups

- rules and principles

# Rule I

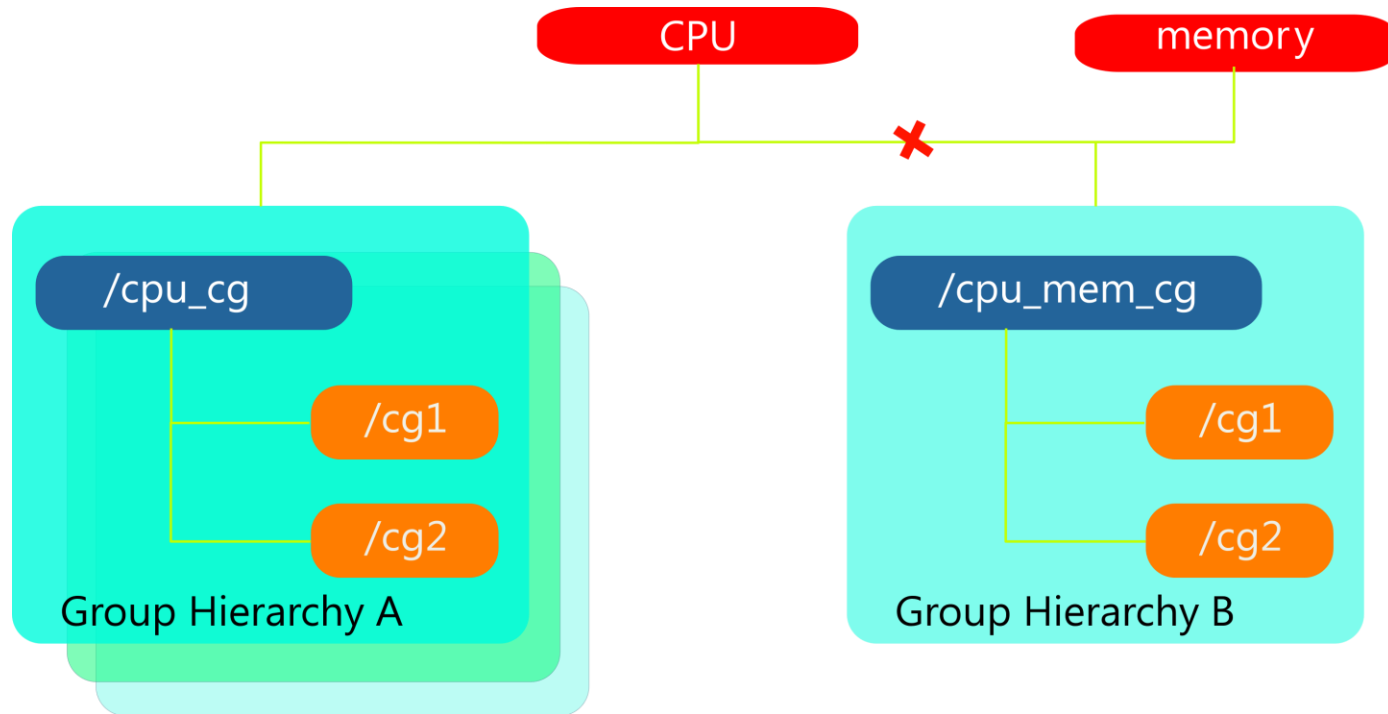
“A single hierarchy can have one or more subsystems attached to it.”





# Rule II


Any single subsystem (such as cpu) cannot be attached to more than one hierarchy if one of those hierarchies has a different subsystem attached to it already.



Note: A single subsystem can be attached to more than one hierarchies as long as they do not have any other subsystem attached already attached

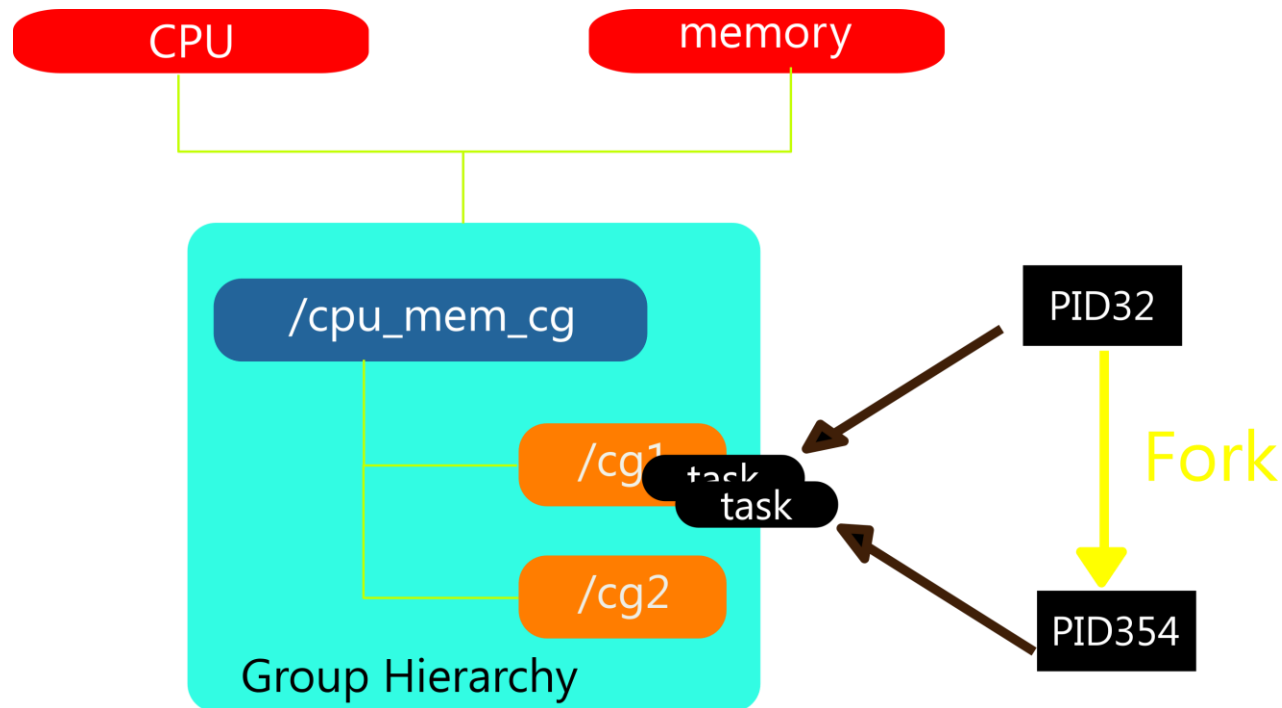
# Rule III

- ❑ New hierarchy → all tasks belongs to root cgroup (recall process organization)
- ❑ A single task can belong to distinct cgroups of distinct hierarchies
  - ❑ Recall “sharing leaves”
- ❑ A single task can not belong in distinct cgroups of the same hierarchy



# Rule IV

- ❑ Process fork → child task
- ❑ Parent membership inheritance (recall process fork inheritance)



# Getting our hands dirty!

---

`cgroups -$ ./tutorial`

(1) Introduction

(2) Using cgroups

> 2

Using cgroups

- | - Files and cgroup VFS
- | - Hands on
- | - The cpu subsystem

> \_


# Files and cgroup VFS

```
cgroups -$ ls /sys/fs/cgroup
blkio cpu cpuacct cpuset devices freezer hugetlb memory
perf_event system
```

```
cgroups -$ ls /sys/fs/cgroup/cpu
cgroup.clone_children cgroup.procs cpu.cfs_period_us cpu.shares
notify_on_release tasks cgroup.event_control
cgroup.sane_behavior cpu.cfs_quota_us cpu.stat release_agent
myCgroup
```

## Hint:

- tasks → list of attached tasks by pid
- cgroup.procs → list of thread group IDs in the cgroup
- notify\_on\_release → flag, “run the release agent on exit?”
- release\_agent → path to use for release notifications
- **Other files** → depends on the policy expression model



# Files and cgroup VFS

---

```
cgroups -$ ls /proc  
... cgroups ...
```

→ available cgroup hierarchies

```
cgroups -$ ls /proc/764  
... cgroup ...
```

→ what this task is attached to?  
path relative to the cgroup file system

# Hands on

Temporary fs

Device type

1. `mount -t tmpfs cgroup_root /sys/fs/cgroup`
2. `mkdir /sys/fs/cgroup/cpuset`
3. `mount -t cgroup -ocpuset cpuset /sys/fs/cgroup/cpuset`
4. Create the new cgroup by doing `mkdir`'s and `write`'s (or `echo`'s) the `/sys/fs/cgroup/cpuset` virtual file system.
5. Start a task that will be the "founding father" of the new job.
6. Attach that task to the new cgroup by writing its PID to the `/sys/fs/cgroup/cpuset/tasks` file for that cgroup.
7. `fork`, `exec` or clone the job tasks from this founding father task.

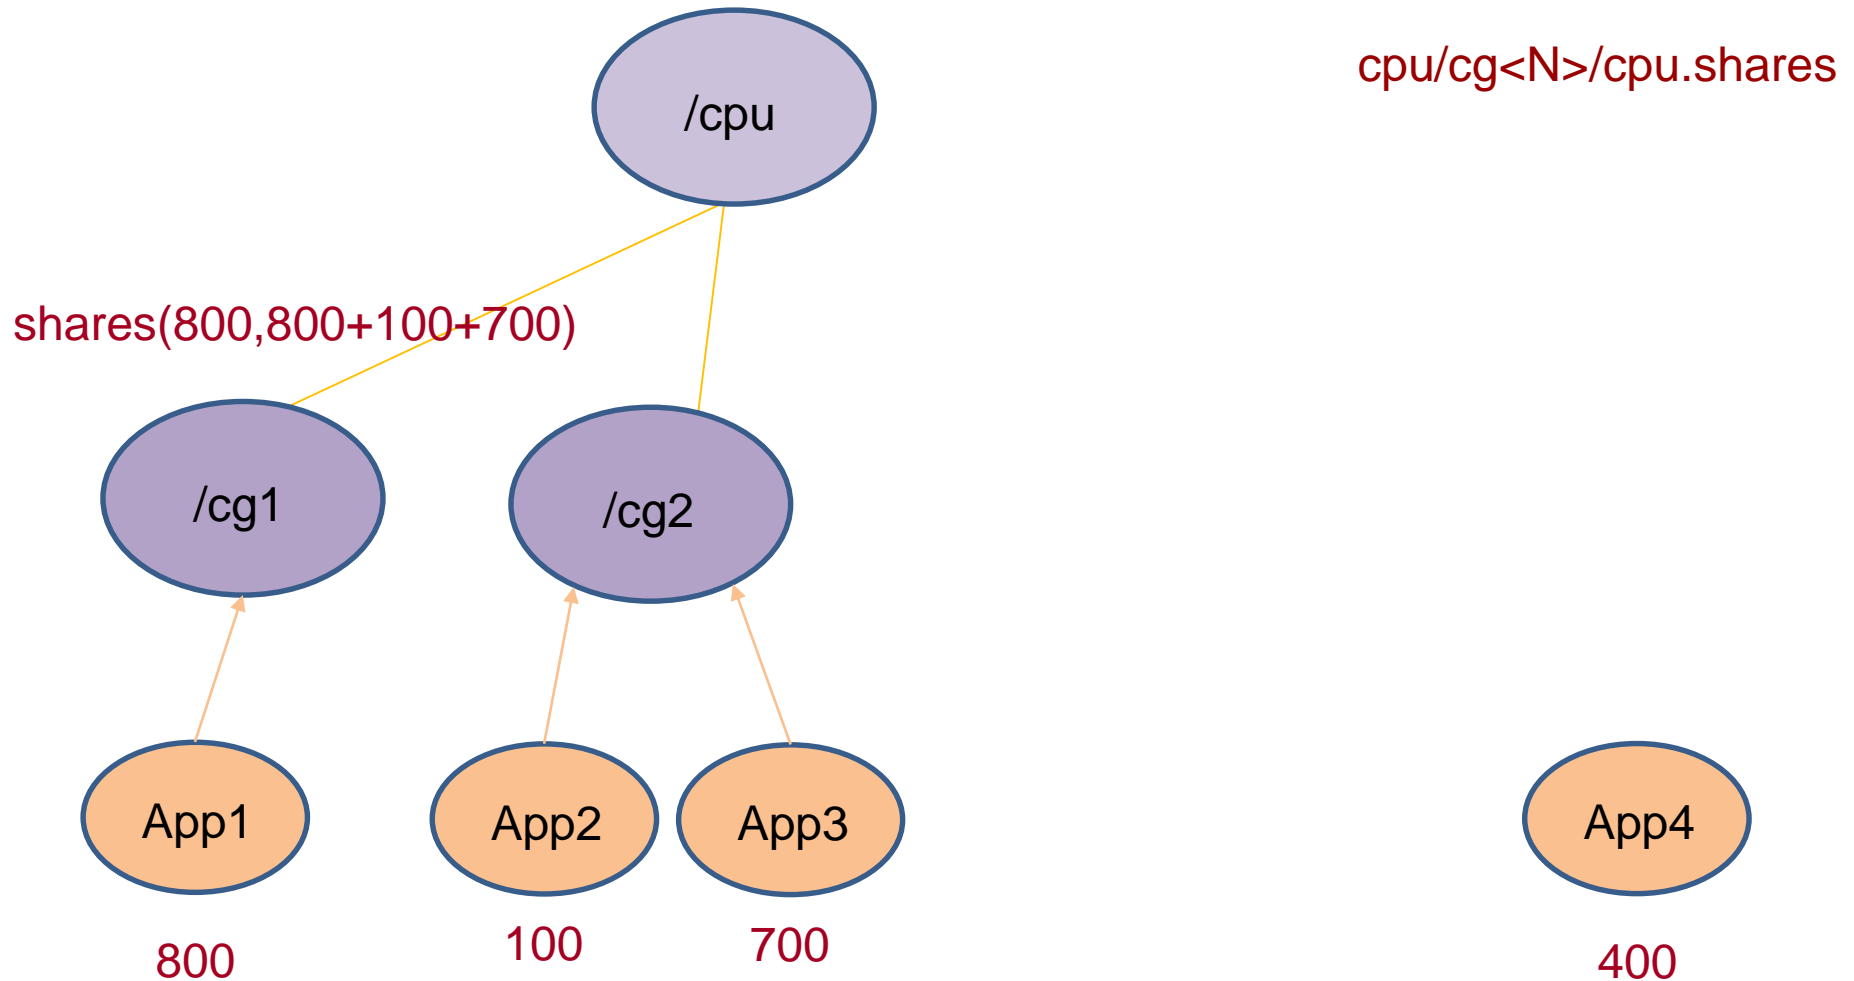
# cpu subsystem

---

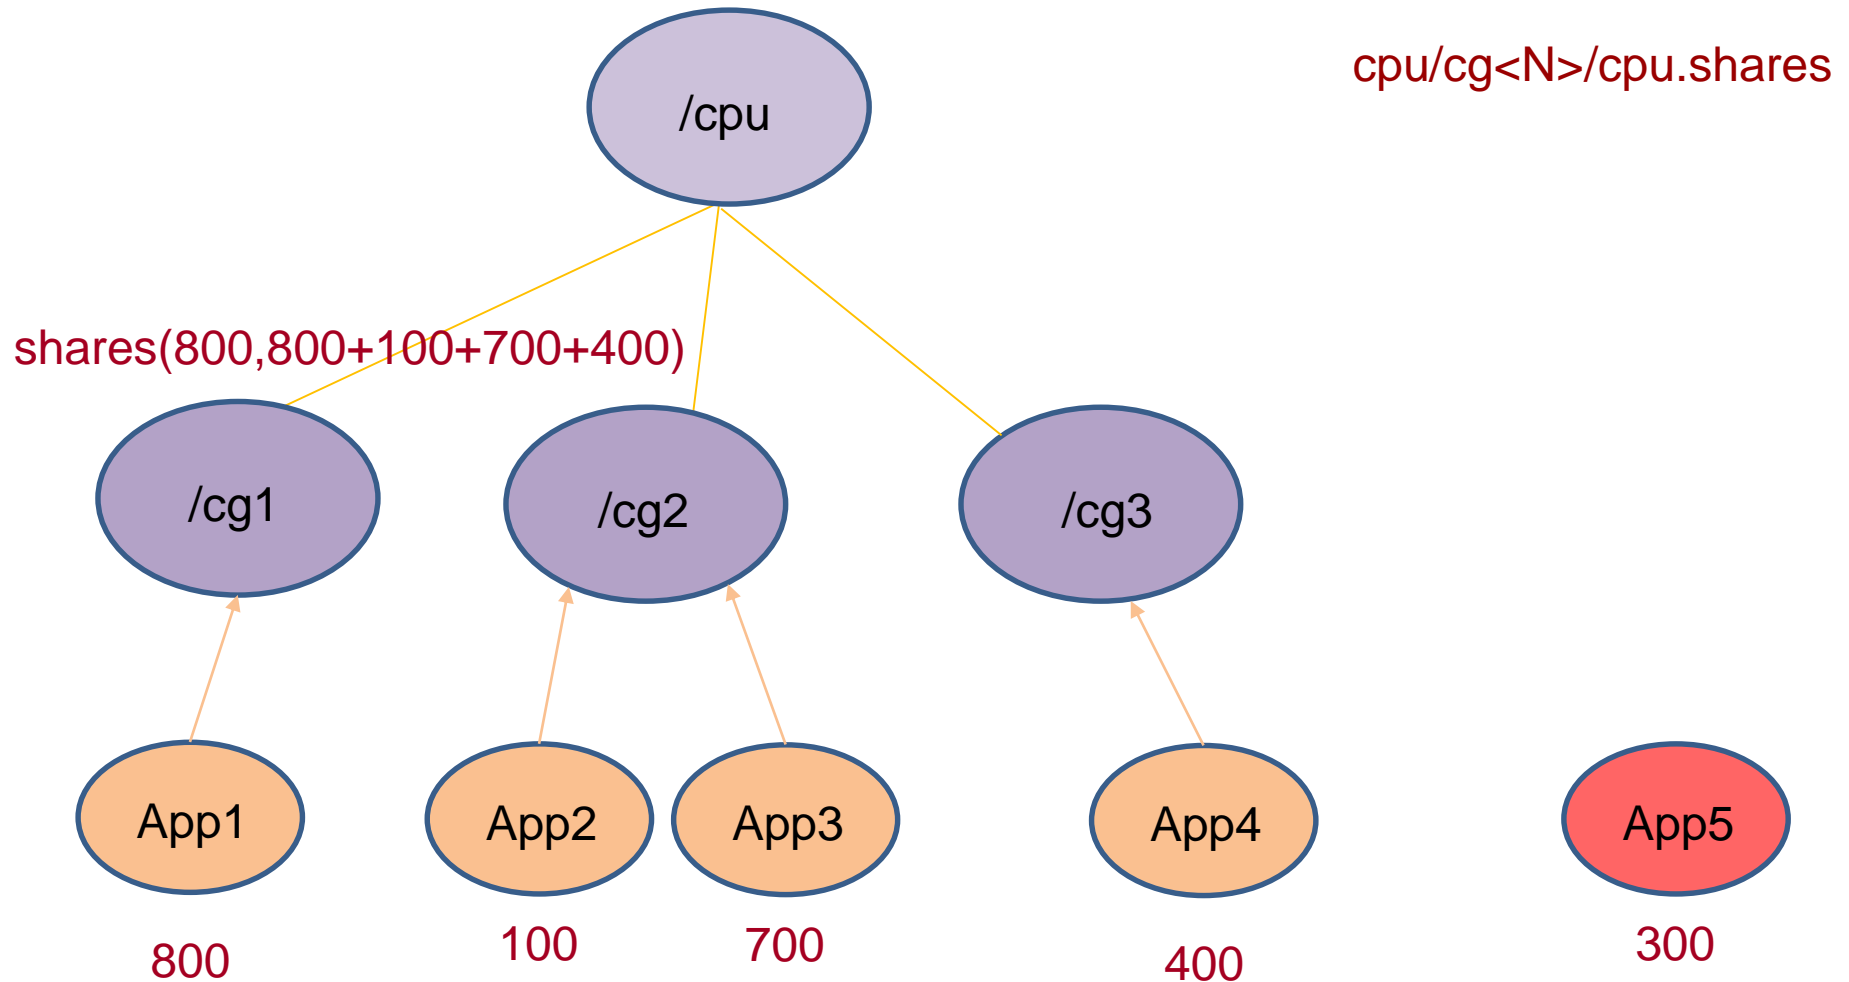
1. **Completely Fair Scheduler (CFS)** — a proportional share scheduler which divides the CPU time (CPU bandwidth) proportionately between groups of tasks (cgroups) depending on the priority/weight of the task or shares assigned to cgroups.
2. **Real-Time scheduler (RT)** — a task scheduler that provides a way to specify the amount of CPU time that real-time tasks can use.  
e.g.: All tasks in a cgroup are allowed to run 0.1 seconds in every 1 second



# cpu subsystem, CFS



# cpu subsystem, CFS



# Άσκηση

## 1. Δαίμων cgmond

## 2. cgmon-policy

- Input: policy:<application name>:cpu:<value>
- Output: score:<float>  
set\_limit:<application name>:cpu.shares:<value>

## 3. cgmon-limit

1. Δημιουργία cgroup για μια νέα εφαρμογή:
  - create:<monitor>:cpu:<application name>
2. Κατάργηση του cgroup μιας εφαρμογής που έχει τερματίσει:
  - remove:<monitor>:cpu:<application name>
3. Εγγραφή μιας διεργασίας στο cgroup μιας εφαρμογής:
  - add:<monitor>:cpu:<application name>:<process id>
4. Ρύθμιση της τιμής cpu.shares για το cgroup μιας εφαρμογής:  
set\_limit:<monitor>:cpu:<applicationname>:cpu.shares:<value>

# For extra info you may refer to ...

---

1. <http://lxr.free-electrons.com/source/kernel/cgroup.c>
2. [https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/6/html/Resource\\_Management\\_Guide/ch01.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Resource_Management_Guide/ch01.html)
3. <http://stackoverflow.com/>